

```

In [4]: import pandas as pd
import numpy as np
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score, C

df = pd.read_csv("C:\\Users\\hp\\Downloads\\WA_Fn-UseC_-HR-Employee-Attrition.csv")

df.drop(['EmployeeNumber', 'Over18', 'StandardHours', 'EmployeeCount'], axis=1, inplace=

# Encode categorical variables
le = LabelEncoder()
for column in df.select_dtypes(include='object').columns:
    df[column] = le.fit_transform(df[column])

# Separate features and target
X = df.drop('Attrition', axis=1)
y = df['Attrition']

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=

# Feature scaling
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Logistic Regression
log_model = LogisticRegression()
log_model.fit(X_train, y_train)
log_pred = log_model.predict(X_test)

# Decision Tree
tree_model = DecisionTreeClassifier()
tree_model.fit(X_train, y_train)
tree_pred = tree_model.predict(X_test)

# Random Forest
rf_model = RandomForestClassifier()
rf_model.fit(X_train, y_train)
rf_pred = rf_model.predict(X_test)

# Evaluation
print("Logistic Regression:\n", classification_report(y_test, log_pred))
print("Decision Tree:\n", classification_report(y_test, tree_pred))
print("Random Forest:\n", classification_report(y_test, rf_pred))

```

## Logistic Regression:

	precision	recall	f1-score	support
0	0.91	0.98	0.94	255
1	0.70	0.36	0.47	39
accuracy			0.89	294
macro avg	0.80	0.67	0.71	294
weighted avg	0.88	0.89	0.88	294

## Decision Tree:

	precision	recall	f1-score	support
0	0.87	0.89	0.88	255
1	0.15	0.13	0.14	39
accuracy			0.79	294
macro avg	0.51	0.51	0.51	294
weighted avg	0.77	0.79	0.78	294

## Random Forest:

	precision	recall	f1-score	support
0	0.88	1.00	0.93	255
1	0.80	0.10	0.18	39
accuracy			0.88	294
macro avg	0.84	0.55	0.56	294
weighted avg	0.87	0.88	0.83	294

```
In [4]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, roc_curve, auc
import matplotlib.pyplot as plt

df = pd.read_csv("C:\\Users\\hp\\Downloads\\WA_Fn-UseC_-HR-Employee-Attrition.csv")

# Encode Target Variable (Attrition)
df['Attrition'] = df['Attrition'].apply(lambda x: 1 if x == 'Yes' else 0)

# Encode Categorical Columns
le = LabelEncoder()
for col in df.select_dtypes(include='object').columns:
    if col != 'Attrition':
        df[col] = le.fit_transform(df[col])

# Split Features and Target
X = df.drop('Attrition', axis=1)
y = df['Attrition']

# Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=

# Train Logistic Regression Model
model = LogisticRegression(max_iter=1000)
```

```

model.fit(X_train, y_train)
y_pred = model.predict(X_test)
y_proba = model.predict_proba(X_test)[: , 1]

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=["Stay", "Left"])
disp.plot(cmap=plt.cm.Blues)
plt.title("Confusion Matrix")
plt.show()

# ROC Curve
fpr, tpr, _ = roc_curve(y_test, y_proba)
roc_auc = auc(fpr, tpr)
plt.plot(fpr, tpr, label=f"ROC Curve (AUC = {roc_auc:.2f})", color="darkorange")
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Receiver Operating Characteristic (ROC) Curve")
plt.legend(loc="lower right")
plt.grid(True)
plt.show()

```

C:\ProgramData\anaconda3\Lib\site-packages\sklearn\linear\_model\\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.


Increase the number of iterations (max\_iter) or scale the data as shown in:


<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n\_iter\_i = \_check\_optimize\_result(

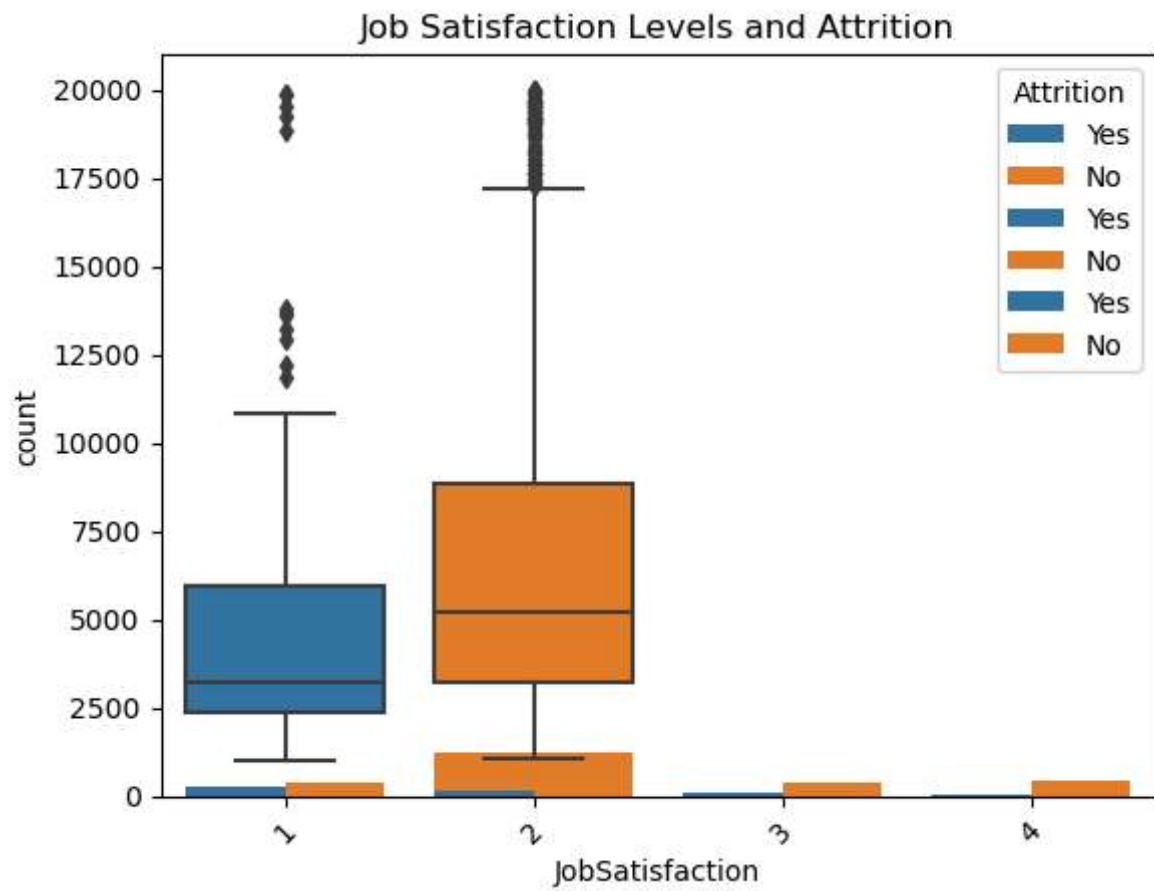




```
In [3]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
df = df = pd.read_csv("C:\\Users\\hp\\Downloads\\WA_Fn-UseC_-HR-Employee-Attrition.csv")

sns.countplot(data=df, x='Attrition')
plt.title("Distribution of Employee Attrition")
plt.xlabel("Attrition (1 = Yes, 0 = No)")
plt.ylabel("Number of Employees")
sns.countplot(data=df, x='Department', hue='Attrition')
plt.title("Attrition Count by Department")
plt.xticks(rotation=45)
sns.boxplot(x='Attrition', y='MonthlyIncome', data=df)
plt.title("Monthly Income and Attrition")
sns.countplot(data=df, x='OverTime', hue='Attrition')
plt.title("Attrition Based on Overtime")
sns.histplot(df[df['Attrition'] == 1]['Age'], kde=True, bins=10)
plt.title("Age Distribution of Employees Who Left")
sns.countplot(data=df, x='JobSatisfaction', hue='Attrition')
plt.title("Job Satisfaction Levels and Attrition")
```

```
Out[3]: Text(0.5, 1.0, 'Job Satisfaction Levels and Attrition')
```



In [ ]: